

From Tweets to Wire Copy: Transferring BEND Maneuver Detection to Long-Form State Media with a Large Language Model Coder

Cynthia Neil

Georgia Institute of Technology, Atlanta GA 30332, USA
cneil7@gatech.edu

Abstract. The BEND framework catalogues the maneuvers by which influence operations act on a narrative and on a community, but its operational definitions and automated detectors were built on short-form social media, where the signal lies in hashtags, emoji, and the retweet graph. Whether it transfers to long-form state-media prose—and whether a large language model (LLM) can code that prose as reliably as a trained human—is untested. We apply Blane’s operational BEND definitions to 150 English-language TASS, RT, and Sputnik articles on Ukraine (January 2021–23 February 2022) retrieved from the Internet Archive. An annotation manual doubling as a structured LLM prompt was applied by one model to all 150 articles and, blind, by two human analysts to a 30-article subset. Initial agreement was low for every pairing, humans included (maneuver κ : human–human 0.23), and the disagreement concentrated almost entirely on one catch-all category. An explicit decision rule for that category, with both humans re-coding blind, raised and converged every maneuver pairing (human–human 0.30; human–LLM 0.32 and 0.33), while thread agreement reached 0.49 and 0.55 against a human–human 0.38. Human–LLM agreement thus meets or exceeds human–human on both coded dimensions, and exceeds published cue-threshold benchmarks several-fold on the maneuver each detects best. Reliability in maneuver coding is a property of the instrument rather than of the coder, and is engineerable.

Keywords: Social cybersecurity · BEND framework · Influence operations · LLM annotation · Inter-coder reliability · Content analysis

1 Introduction

Between spring 2021 and Russia’s full-scale invasion of Ukraine on 24 February 2022, the Kremlin’s English-language state media prepared the informational ground for war. TASS, RT, and Sputnik did not merely report: they performed recurring rhetorical maneuvers—recasting NATO enlargement as encirclement, reframing Ukrainian self-defense as aggression, and escalating toward a “genocide” pretext—aimed at Western policy and elite audiences.

Social cybersecurity offers an operationalized vocabulary for this behavior: Carley’s BEND framework catalogues sixteen influence maneuvers along two

axes [2], which Blane refined into coded operational definitions with cue-based automated detectors [1]. Both were built on *short-form* social media, where detection leans on hashtags, emoji, capitalization, and the reply/retweet graph. Yet much consequential influence content is long-form prose—state wire copy, syndicated commentary, “pink slime” local-news imitations, and the machine-targeted article farms now built to shape model training corpora. Whether BEND’s narrative-axis maneuvers remain *reliably* codable once the short-form cues disappear is open and practically consequential.

A second question rides on the first: instruction-following LLMs increasingly substitute for human annotators in computational social science [5, 16], subject to a standing caution that such substitution be validated rather than assumed [8]. Maneuver coding is hard, interpretive, and low-agreement—precisely the kind of task on which automation claims deserve scrutiny.

This pilot addresses both. We report the first reliability test, to our knowledge, of BEND narrative-axis coding on long-form state-media prose, measured against a human–human baseline rather than a single-coder gold standard; a diagnosis-and-repair result in which the unreliability concentrates in one catch-all category and an explicit priority rule targeting it converges all three coder pairings, with both humans re-coding blind so the gain is measured on the instrument rather than assumed from the model; and evidence that an LLM functions as a human-equivalent coder for this instrument, exceeding the cue-threshold benchmark reported for the same framework on short-form text. Detection is the deliverable throughout; response effectiveness is out of scope.

2 Background

BEND organizes influence maneuvers along two axes: eight *narrative* maneuvers acting on the message (Engage, Explain, Excite, Enhance; Dismiss, Distort, Dismay, Distract) and eight *community* maneuvers acting on the social graph (Back, Build, Bridge, Boost; Neglect, Narrow, Neutralize, Nuke) [2]. Blane supplied the reference operational definitions this study adopts, with NetMapper cue-based detection validated on large Twitter corpora [1]. Her reliability findings are our benchmark, and they are sobering: human–human agreement among lightly-to-moderately trained annotators was “generally weak, at or below” $\kappa = 0.35$; highly trained annotators ranged from -0.01 to 0.53 ; and her automated detector peaked at $\kappa \approx 0.15$ on *dismay*, with most maneuvers near zero or negative.

The framework remains under active extension. Williams et al. carry BEND to webgraphs, proposing metrics for community maneuvers in SEO-manipulated link structures [15]; Hicks operationalizes it at the population and effects levels and frames effects-based detection as *complementary* to message-level content methods [6]. This work sits on that complementary side and extends it along a dimension neither addresses: text length and register.

Separately, a growing literature reports that instruction-following LLMs rival or exceed trained annotators on text-labeling tasks [5, 16], with the caution that automated annotation requires human benchmarking before its labels are

trusted [8]. This study validates its LLM coder against two human analysts, reporting the human–human floor alongside every human–machine figure.

3 Data and Method

3.1 Corpus

The corpus comprises 150 English-language articles from TASS English, RT, and Sputnik English across three sub-periods—spring 2021 (the “false alarm”), late 2021, and the pre-invasion window of January–23 February 2022—at roughly fifty each ($P1 = 50$, $P2 = 49$, $P3 = 51$) with an even per-period outlet mix. These outlets constitute the Kremlin’s principal English-language *export* messaging layer, aimed at foreign elite rather than domestic Russian audiences, and are documented as such in U.S. government reporting [13, 9]. Every article was retrieved from the Internet Archive Wayback Machine rather than scraped live—avoiding interaction with sanctioned outlets’ live infrastructure and yielding a citable, immutable snapshot per item, with original URL, verified snapshot timestamp, and extracted body text recorded for each. Articles carry a period–outlet–number identifier (e.g. P2-RT-13).

3.2 Coding instrument

The unit of analysis is the whole article, coded with Blane’s operational definitions of the *eight narrative-axis maneuvers*. The eight community maneuvers are treated as out of modality: their cues live in the reply/retweet/follow graph that a standalone article does not carry. Each article receives one primary maneuver, ranked secondary and usually tertiary maneuvers, a narrative thread, supporting evidence quotes, and a confidence rating. Two threads are locked from the public attribution record: T1 (NATO expansion as *casus belli*) and T2 (neo-Nazi / illegitimate-regime framing), with “both” and “other” available. Critically, the instrument is a *single artifact* serving as both the human annotation manual and the structured LLM prompt, so human and machine outputs are directly comparable and any refinement applies identically to both. It was authored as v1 and, in response to the reliability findings reported below, revised to v1.1.

3.3 Classification and validation

All 150 articles were classified by Claude Opus 4.8 using the manual as a cached system prompt, with adaptive thinking and schema-constrained structured output returning one JSON object per article. Only neutral metadata (identifier, outlet, date, headline) and body text were sent; the manifest’s keyword-derived thread hypotheses were withheld to avoid biasing the coder. The run completed with zero errors.

A 30-article validation subset was drawn, stratified by the model’s primary-maneuver label across the four categories that populate the corpus and balanced

across sub-period and outlet. Two human coders—the author and an independent peer analyst—coded it blind to the model and to one another, using the same manual. Agreement was quantified with Cohen’s κ [3] across three pairings (human–human, and each human against the LLM) on two dimensions (primary maneuver and narrative thread), computed both as a multi-class statistic and, matching Blane’s construct, as a per-maneuver binary statistic in which presence means primary or secondary. The subset was coded under v1 and, after refinement, re-coded from scratch under v1.1 by all three coders, yielding a controlled before/after comparison.

4 Results

4.1 Reliability under the initial instrument

Under v1 (Table 1, left columns), *thread* agreement was moderate in Landis and Koch’s terms [7] and—notably—the human–LLM pairings matched human–human: 0.44 (author–peer), 0.41 (author–LLM), 0.47 (peer–LLM), with fourteen of thirty articles drawing full three-way agreement. *Primary maneuver* agreement was much lower, and low *for the humans too*: 0.23, 0.29, and 0.05 respectively. The LLM agreed with the author slightly better than the two humans agreed with each other. A result of this shape is easy to misread as machine failure; the human–human floor shows it is not.

4.2 Where the disagreement lives

The disagreement is not diffuse. Across the human–LLM pairings the dominant confusions are Distort $\leftrightarrow$ Explain (11 instances), Dismay $\leftrightarrow$ Distort (10), and Dismiss $\leftrightarrow$ Distort (7): every high-frequency confusion pivots on **Distort**. The two humans coded Distort as primary on 19 and 21 of 30 articles respectively, against the model’s 8, and where a human coded Distort the model agreed only about 35% of the time.

The interpretation matters. Where slanted framing is pervasive, “Distort” invites use as a default for any tendentious article, absorbing cases whose dominant function is actually emotional (Dismay), dismissive (Dismiss), or expository (Explain). Because *the two humans disagree on the same boundary*, this is a property of the instrument, not of the model.

Thread disagreement was parallel: T1 $\leftrightarrow$ other (8) and T2 $\leftrightarrow$ other (6) dominate. Where the author coded “other”—typically sub-narratives of Western hypocrisy or U.S. incompetence—the peer and the model more often folded the article back into T1: a convention gap never pinned before coding, not a perceptual difference.

4.3 Refinement and re-test

Two targeted revisions produced v1.1. First, a **Distort decision rule**: a priority test treating factual slant as the *substrate* of this corpus rather than its

label, requiring the coder to identify the dominant rhetorical *function* as primary (Dismay $\rightarrow$ Dismiss $\rightarrow$ Distract $\rightarrow$ Explain $\rightarrow$ Distort, first match wins) and to record falsity as a secondary Distort against a citable external anchor. Distort as primary now requires both dominance and an anchorable factual claim. Second, a **thread convention**: the two threads are *casus belli* pillars, so an article must actually argue a war-justification to earn a thread code; free-standing Western-hypocrisy material is coded “other” rather than forced into T1. A third change instructed all coders to record secondary and tertiary maneuvers, equalizing label depth.

Results of the re-coding are in Table 1.

Table 1. Cohen’s κ before and after instrument refinement, $N = 30$. Both human coders re-coded blind under v1.1.

Pairing	Primary maneuver		Narrative thread	
	v1	v1.1	v1	v1.1
Author–peer (human–human)	0.232	0.302	0.437	0.378
Author–LLM	0.292	0.324	0.408	0.492
Peer–LLM	0.054	0.332	0.470	0.552

Three findings stand out. **The maneuver pairings converged:** a v1 spread of 0.05–0.29 collapsed to a tight 0.30–0.33, and the near-chance peer–LLM pairing rose six-fold, confirming that its original value reflected inconsistent use of the catch-all rather than a perceptual gap. **Human–LLM agreement now meets or exceeds human–human on both dimensions**—maneuver 0.324 and 0.332 against 0.302, thread 0.492 and 0.552 against 0.378—so the model agrees with each human at least as well as the two humans agree with each other, and without treating either human as ground truth. **The mechanism is visible and intended:** both humans sharply de-defaulted Distort, from 19 and 21 of 30 under v1 to 1 and 7 under v1.1 against the model’s 2, redistributing those articles to the emotional and deflection functions the rule prioritizes—of the author’s 19 v1 Distorts, 10 became Dismay and 5 Distract.

One effect ran the other way and we report it as-is: **human–human thread agreement fell, 0.437 to 0.378.** The richer convention introduced genuine judgment at the “both”/“other” boundary, where the humans diverged—the author used “both” four times, the peer not at all. The convention improved machine alignment, since the model applies it consistently, while exposing that the “both” call still needs a tighter human-facing rule: the one boundary the refinement did not close.

4.4 Comparison with cue-based automation

Blane’s cue-based detector peaked at $\kappa \approx 0.15$ on *dismay* against human coders (Section 2). On the same per-maneuver binary construct under v1.1, human–

LLM κ on Dismay here is 0.500 and 0.533—the very maneuver on which her automation performed best, exceeded by roughly three and a half times. Notably, human–human agreement on Dismay in our data is only 0.133: each human agrees with the model considerably more than with the other.

Two comparisons follow. First, **the reliability profile transfers**: our human–human agreement—multi-class 0.30, and under v1 a per-maneuver maximum of 0.37—sits inside Blane’s lightly-to-moderately-trained band and below her highly-trained ceiling, so BEND coding is about as hard in long-form prose as in short-form tweets. Second, **an LLM reading full articles is a markedly better automated detector than a cue-threshold pipeline**, closing a human–machine gap that was wide on Twitter.

These comparisons are indicative rather than exact: Blane’s estimates rest on 1,300–2,999 items against our 30, and low-prevalence maneuvers yield unstable κ , so only Distort, Dismiss, and Dismay should be read.

5 External Validity

If the codes are meaningful rather than idiosyncratic, they should recover structure an independent source also identifies. Re-coding all 150 articles under v1.1 yields primary maneuvers Explain 57, Dismiss 41, Dismay 37, Distort 7, Enhance 4, Distract 2, Excite 0. Distort as a *primary* label is now rare, but as a *property* it is pervasive—primary or secondary in 95 of 150 articles, exactly what the decision rule intends. Excite, meanwhile, is empirically empty, corroborating the export-audience scope.

Ninety-three of 150 articles (62%) advance one of the two threads named in the U.S. attribution record [11, 12]; the remaining 57 code as “other.” Each cluster carries a distinct primary-maneuver signature: the NATO/security grievance *argues* (Explain 30 of 43), the regime-delegitimization narrative *frightens* (Dismay 29 of 50), and the residual cluster—a coherent Western-illegitimacy strand rather than noise—*deflects* (Dismiss 30 of 57). The clusters differ not merely in topic but in method, a separation the unrefined instrument could not show because inflated Distort counts sat on top of all three. The temporal movement is likewise carried by thread rather than maneuver: the regime cluster nearly doubles between the first and third sub-periods (13 to 24) as the messaging turns to pretext, matching the escalation the U.S. government attributed contemporaneously [13, 14].

6 Discussion and Limitations

What transfers, and what needed modification. The narrative-axis maneuvers are recoverable from long-form text at coarse grain, and after refinement human–LLM agreement reaches or exceeds human–human on both dimensions. Three modifications were required: the community maneuvers do not transfer to single-article text and are properly excluded; Blane’s short-form cue indicators

do not survive, the productive long-form substitutes being scare-quoting of contested terms (‘invasion’, ‘Nazis’, ‘genocide’), attribution and sourcing patterns, and within-article juxtaposition; and the Distort definition was under-specified for prose in which slant is pervasive.

The general lesson. Reliability here was a property of the *instrument*, not of the coders. The diagnostic sequence—measure a human–human floor, locate the confusions, find that they pivot on one over-broad category, write an explicit priority rule, then re-code blind with every coder—is portable to any maneuver-coding effort, with analogous catch-alls likely in adjacent taxonomies—DISARM’s overlapping “five Ds” [4] an obvious place to look.

Limitations. The validation subset is small ($N = 30$), so confidence intervals are wide, on the order of ± 0.2 ; these are pilot-grade estimates. The author served as one of the two human coders, so that pairing is not fully independent; the peer–LLM pairing is the cleaner comparison and also the one that improved most. Reliability was established on the subset, while the remaining 120 articles carry LLM codes whose agreement with human judgment is inferred rather than measured. Coder calibration on the new rule varied, one coder applying it more aggressively than the model. The design covers one language, one actor, one model, and one time window. Finally, reliability is not validity: agreement establishes consistency, not ground truth, which is why the corpus-level codes are anchored to independent government attributions in Section 5.

What detection is for. Detection is a means. A parallel strand of this work develops a democratically-constrained response framework, on the premise that the maneuvers coded here are in most cases protected speech a state may neither suppress nor credibly mirror—which confines the response space to attribution, truth-anchoring, and inoculation [10]. A reliable, auditable maneuver label is the input that framework requires, and it is the reason reliability was worth engineering rather than merely reporting.

7 Conclusion

An operationalized short-form influence-maneuver framework transfers to long-form state-media prose at a reliability that is genuinely usable once a single catch-all category is disciplined with an explicit decision rule. At that point a large language model codes the material about as consistently as a trained human analyst—meeting or exceeding human–human agreement on both coded dimensions—and substantially better than the cue-threshold automation the framework’s own author benchmarked. The resulting corpus-level codes recover a three-strand messaging architecture whose escalation tracks the public attribution record. We offer the diagnose-and-repair protocol, rather than the particular numbers, as the transferable contribution: in low-agreement interpretive coding, the instrument is the thing to fix.

Acknowledgements. The author thanks James Bonnell for serving as the independent second coder, blind-coding the validation subset under both versions of

the instrument; the human–human baseline central to this paper would not exist without that contribution. The author has no competing interests to declare.

References

1. Blane, J.T.: Social-Cyber Maneuvers for Analyzing Online Influence Operations. Ph.D. thesis, Carnegie Mellon University (2023). CMU-S3D-23-102. <https://doi.org/10.1184/R1/22825112>
2. Carley, K.M.: Social cybersecurity: an emerging science. *Computational and Mathematical Organization Theory* **26**(4), 365–381 (2020). <https://doi.org/10.1007/s10588-020-09322-9>
3. Cohen, J.: A coefficient of agreement for nominal scales. *Educational and Psychological Measurement* **20**(1), 37–46 (1960). <https://doi.org/10.1177/001316446002000104>
4. DISARM Foundation: DISARM Frameworks. GitHub repository (accessed 11 August 2026). <https://github.com/DISARMFoundation/DISARMframeworks>
5. Gilardi, F., Alizadeh, M., Kubli, M.: ChatGPT outperforms crowd workers for text-annotation tasks. *PNAS* **120**(30), e2305016120 (2023). <https://doi.org/10.1073/pnas.2305016120>
6. Hicks, M.: AI-Enabled Social-Cyber Maneuver Detection and Creation. Ph.D. thesis, Carnegie Mellon University (2025). <https://doi.org/10.1184/R1/29316026>
7. Landis, J.R., Koch, G.G.: The measurement of observer agreement for categorical data. *Biometrics* **33**(1), 159–174 (1977). <https://doi.org/10.2307/2529310>
8. Pangakis, N., Wolken, S., Fasching, N.: Automated annotation with generative AI requires validation. *arXiv:2306.00176* (2023). <https://arxiv.org/abs/2306.00176>
9. Paul, C., Matthews, M.: The Russian “Firehose of Falsehood” Propaganda Model. RAND Corporation, Santa Monica (2016). <https://www.rand.org/pubs/perspectives/PE198.html>
10. Roozenbeek, J., van der Linden, S.: Fake news game confers psychological resistance against online misinformation. *Palgrave Communications* **5**, 65 (2019). <https://doi.org/10.1057/s41599-019-0279-9>
11. U.S. Department of State: Fact vs. Fiction: Russian Disinformation on Ukraine. U.S. Mission Geneva, 20 January 2022. <https://geneva.usmission.gov/2022/01/20/fact-vs-fiction-russian-disinformation-on-ukraine/>
12. U.S. Department of State: Russia’s Top Five Persistent Disinformation Narratives, 20 January 2022. <https://2021-2025.state.gov/russias-top-five-persistent-disinformation-narratives/>
13. U.S. Department of State, Global Engagement Center: Kremlin-Funded Media: RT and Sputnik’s Role in Russia’s Disinformation and Propaganda Ecosystem, January 2022. https://2021-2025.state.gov/wp-content/uploads/2022/01/Kremlin-Funded-Media_January_update-19.pdf
14. U.S. Department of State: Disinformation Roulette: The Kremlin’s Year of Lies to Justify an Unjustifiable War, February 2023. <https://2021-2025.state.gov/disarming-disinformation/disinformation-roulette-the-kremlins-year-of-lies-to-justify-an-unjustifiable-war/>
15. Williams, E.M., Carragher, P., Herdrich, K., Prakarsa, L., Carley, K.M.: Extending the BEND framework to webgraphs. In: *Social, Cultural, and Behavioral Modeling (SBP-BRIMS 2025)*. LNCS. Springer (2025). https://doi.org/10.1007/978-3-032-07715-8_8
16. Ziemis, C., Held, W., Shaikh, O., Chen, J., Zhang, Z., Yang, D.: Can large language models transform computational social science? *Computational Linguistics* **50**(1), 237–291 (2024). https://doi.org/10.1162/coli_a_00502